Computer Networks

Lab #4 - Using the HTTP Protocol to Download Digital Objects from a Server

Summary

rc2018

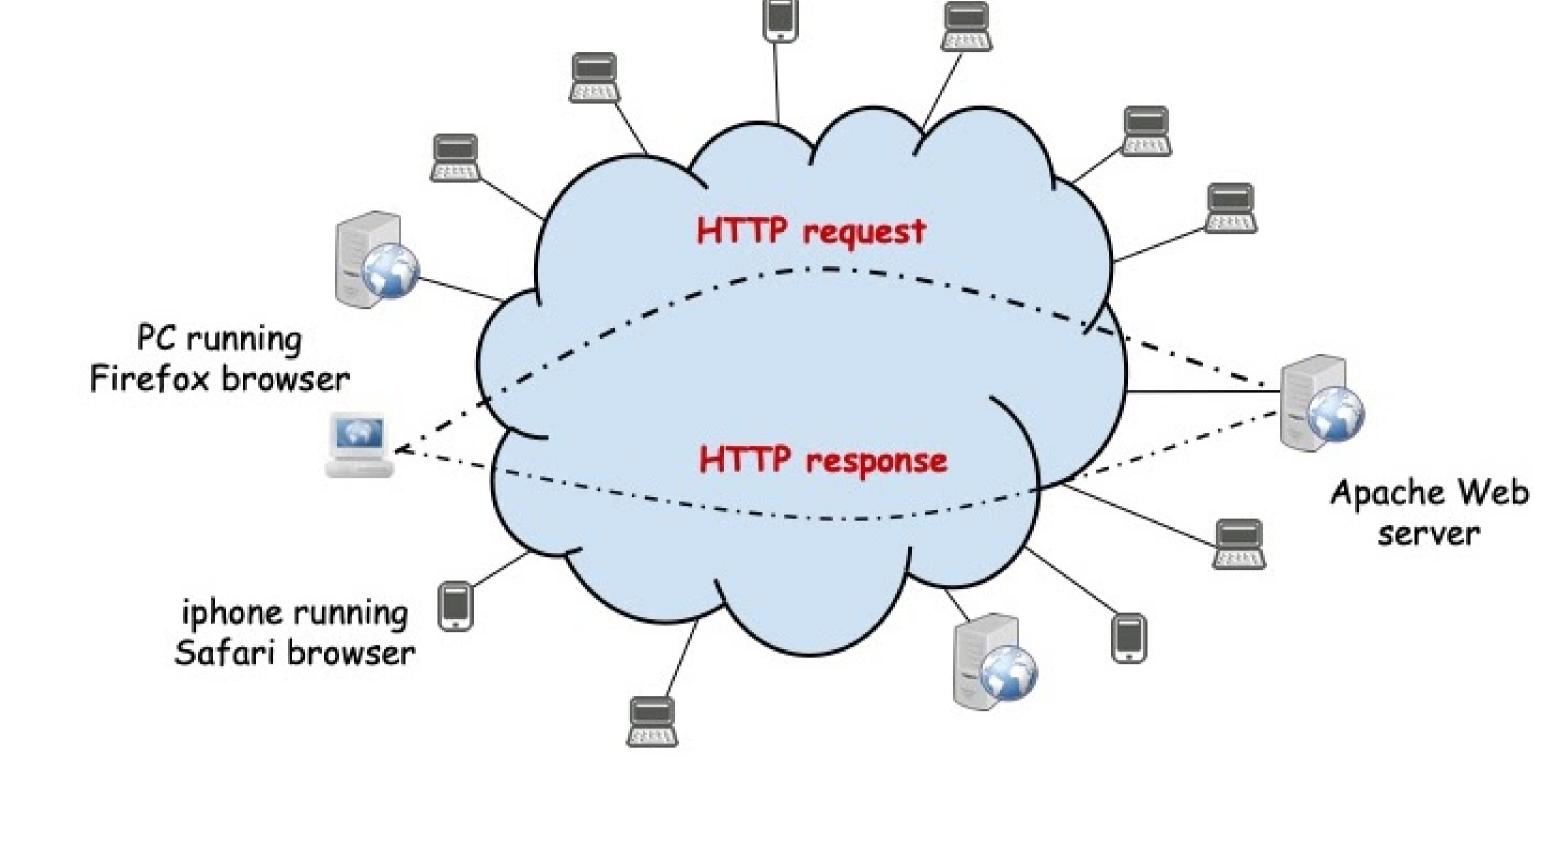
- HTTP Protocol
- Examples and utilities HTTP gets using ranges File transfer over HTTP with and without ranges

HTTP Protocol

- HTTP is a client / server protocol
- The server expects the client to open a TCP connection to its port (by default the port 80) ■ In version 1.0 of the protocol, each request / reply HTTP transaction uses a different TCP connection
- In version 1.1 several HTTP transactions can share the same TCP connection

Hyper Text Transfer Protocol

Overview - Labs - More 🗸



HTTP Messages

```
GET /index HTTP/1.0 CRLF
         Host: www.some-example.com CRLF
                                                     HTTP Request
         User-Agent: Mozilla/40.01 CRLF
                                                        message
         CRLF
                 HTTP/1.1 200 OK CRLF
                 Date: Mon, 21 Feb 2015 17:51:37 GMT CRLF
                 Server: Apache/2.4.10 (Debian) CRLF
                 Last-Modified: Wed, 17 Aug 2014 17:46:43 GMT CRLF
                 ETag: "1267-5272fc1726880" CRLF
HTTP Reply
                Accept-Ranges: bytes CRLF
  message
                 Content-Length: 4768 CRLF
                 Vary: Accept-Encoding CRLF
                 Connection: close CRLF
                 Content-Type: text/html CRLF
                 CRLF
                                                                                       Tempo
Nota: CRLF representa os códigos dos caracteres de controlo carriage return e line feed
```

HTTP request header-fields

HTTP requests made "by hand"

telnet asc.di.fct.unl.pt 80

.... analyze the result.

Or the following one:

long list.

Request Header-fields

There are many different request header-fields that the client can send to the sever into the request message header. Below you will find some of a very

Header-fields Exemplo

User-Agent	User-Agent: Mozilla/40.0
Accept-Charset	Accept-Charset: utf-8
Accept-Encoding	Accept-Encoding: gzip
Accept-Language	Accept-Language: en-UK
If-Modified-Since	If-Modified-Since: Tue, 02 Feb 2016 14:25:41 GM
lf-Match	If-Match: "756154ad8d 3102f1349317f"
Range	Range: bytes=500-999

GET / HTTP/1.0 <return> <return> analyze the result.

header-fields contains several informations, namely the reply object meta data, that are usefull to the client.

By using the following commands, you can access an HTTP server to see its replies. Try the following one:

Or the following one: telnet www.google.com 80 GET / HTTP/1.0 <return> <return>

If your system doesn't have the **telnet** command, you can use instead the **nc** command: nc -c asc.di.fct.unl.pt 80 GET / HTTP/1.0 <return> <return> analyze the result.

nc -c www.google.com 80 GET / HTTP/1.0 <return> <return> analyze the result.

Server

Last-Modified

Header-fields

HTTP reply header-fields

Reply Header-fields

Exemplo

The server sends information to the client by also using the header of the reply message, which is composed of different reply header-fields. These

Server: Apache Last-Modified: Tue, 02 Feb 2016 14:25:41 GMT

Content-Type	Content-Type: text/html; charset=utf-8
Content-Length	Content-Length: 348
Content-Encoding	Content-Encoding: gzip
ETag	ETag: "3d2-52aca46b79fd9"
Accept-Ranges: bytes	Accept-Ranges: bytes
•••••	
code examples	

provides some extra methods to facilitate the development of Java programs build directly on top of the HTTP Protocol. Below you will find several Java source code examples

can only use the class URL to parse an url.

Parsing an URL and opening a TCP connection to the server String url = args[0]; // for example "http://google.com" URL u = new URL(url);// Assuming URL of the form http://server-name/path

There are many other classes available in the same package to develop programs based in the HTTP protocol. However, due to pedagogical reasons, you

Any other requierements of your programs should be implemented by yourself or using the class Http, available in the source code repository, which

• Class <u>URL</u> - allows parsing an url to, for example, get its different components. See file URLparse.java in the <u>source code repository</u>

int port = u.getPort() == -1 ? 80 : u.getPort(); String path = u.getPath() == "" ? "/" : u.getPath(); Socket sock = new Socket(u.getHost(), port); OutputStream out = sock.getOutputStream();

"GET %s HTTP/1.0\r\n"+ "Host: %s\r\n"+ "User-Agent: X-RC2018\r\n\r\n", path, u.getHost()); out.write(request.getBytes());

Composing and sending a request to the server

String request = String.format(

sendFile(request[1], out);

sendsNotSupportedPage(out);

throws IOException {

String page =

// send payload

if(n == -1) break;

http://google.com

http://www.google.com

for(;;) {

} else {

one, among many, that can be used to parse and access URLs:

Parsing the request message header in the server line = Http.readLine(in); String[] request = Http.parseHttpRequest(line); line = Http.readLine(in); while (! line.equals("")) { line = Http.readLine(in); if(request[0].equalsIgnoreCase("GET") && request[1] != "") {

/** * Sends an error message "Not Implemented" private static void sendsNotSupportedPage(OutputStream out)

Example - sending a reply message to the client

"<HTML><BODY>Demo server: request Not Implemented</BODY></HTML>"; int length = page.length(); String header = "HTTP/1.0 501 Not Implemented\r\n"; header += "Date: "+new Date().toString()+"\r\n"; header += "Content-type: text/html\r\n"; header += "Server: "+"X-Server-RC2018"+"\r\n"; header += "Content-Length: "+String.valueOf(length)+"\r\n\r\n"; header += page; out.write(header.getBytes()); Sending the reply header and payload File f = new File(name); long size = f.length();

out.write(buffer, 0, n); **Program demos - HTTP client and server**

Study its code and try to access some urls like for example:

FileInputStream file = new FileInputStream(f);

out.write(header.toString().getBytes());

byte[] buffer = new byte[1024];

int n = file.read(buffer);

header.append("Server: "+"X-Server-RC2018"+"\r\n");

StringBuilder header = new StringBuilder("HTTP/1.0 200 OK\r\n");

header.append("Content-Length: "+String.valueOf(size)+"\r\n\r\n");

http://asc.di.fct.unl.pt http://asc.di.fct.unl.pt/rc Explain the output and understand how it works.

The provided client is able to use the HTTP request / reply protocol to obtain a file and show its content. Modify class HttpClientDemo in a way that it

You can use the browser of your choice and try to access URL http://localhost:8080. The browser will also show the answer of the server.

In the source code repository you will find a naif HTTP client (HttpClientDemo.java) that is able to request an object denoted by the url passed as

For example, if the server is executing in your localhost, and in its current directory there is a file called "Http.java", using the URL http://localhost:8080/Http.java, allows one to see the contents of that file. **Exercise A**

may be used to download files from a HTTP server. Call your class **getFile** for example.

argument. The reply of the server is shown to the user (but it is not parsed, nor interpreted).

In the source code repository there is another class (URLget) which uses the class URL to download an object from an HTTP server. It is shown just for study purposes. Due to pedagogical reasons, you also cann't use it to complete any of your exercises. Range requests

reply headers-fields. To make a partial request, one needs at least two resources: An HTTP server that understands range requests and replies to them

"Host: %s\r\n"+

■ A client using the range request header-field. Examples:

■ Range: bytes=100-199 // requesting the 100 bytes starting with byte 100 // requesting from byte 100 to the end of the object Range: bytes=100-HTTP request made "by hand"

During the (already very long) life of the HTTP protocol, many extensions have been introduced. Most of those extensions introduced new request and

By using the following commands, you can access an HTTP server to see its replies. Try the following one: telnet asc.di.fct.unl.pt 80

<return> analyze the result. Source code example:

"Range: bytes=100-199\r\n"+ "User-Agent: X-RC2018\r\n\r\n", path, u.getHost()); out.write(request.getBytes());

it, you should use the RTFC method (Read The F. Code), which is the only accessible method when no manual is available. To test the actions of the lazy server, launch it in a directory where you also put the file **Earth.jpg** (which size is around 13 Mbytes). You can access it by using the url http://localhost:8080/Earth.jpg with the browser of your choice.

Range requests and replies are specially useful to deal with multimedia information as you will see in the next week's work.

able to fully transfer a file of any size from the HttpLazyServer.java. Since in your future work, you will be directed to use docker hosted web servers, you can also test your client with a different HTTP server. This one is available in a docker container and can be loaded using the command:

In the source code repository you will also find a lazy HTTP server (HttpLazyServer.java) that is able to serve the requested files and partially supports

ranges. In fact, the full support of RFC 7233 ranges is quite complex and extensive. The provided server only supports ranges of the forms shown above.

This server is lazy since it only sends at most MAX_BYTES bytes in each reply. You can find this constant in its source file. Therefore, if you want to know

docker run -it --rm -p 8080:8080 -v \$(PWD):/public danjellz/http-server

You must run this docker in the current directory of the content you want to serve.

Class: FileOutputStream A file can be open in **Append mode**

File f = new File (fileName); Class: File

the method length() can be used to know the **length of a file**. File f = new File (fileName);

long size = f.length(); Class: RandomAccessFile

File f = new File (fileName); RandomAccessFile file = new RandomAccessFile (f, "rw"); file.skipBytes(n); // Attempts to skip over n bytes of input discarding the skipped bytes. file.seek(k); // Sets the file-pointer offset, measured from the beginning of this file, at which the next read or write occurs.

n = file.read(buffer, 0, len); // Reads up to len (or buffer.length) bytes of data from this file into an array of bytes

In the source code repository you will also find a naif HTTP server (HttpServerDemo.java) that is able to serve the requested files from its local file system. For example, if the server is running in the same machine as your browser, you can interact with it using: http://localhost:8080

To facilitate the access to very big or huge objects, HTTP supports partial requests. This feature is inspired from reading direct access files in slices.

GET / HTTP/1.0 <return> Range: bytes=10-20 <return>

OutputStream out = sock.getOutputStream(); String request = String.format("GET %s HTTP/1.0\r\n"+

You can learn about this feature of the HTTP protocol by starting here or looking at RFC 7233.

Exercise - downloading a file using successive range downloads Now, assuming that you already have extended the demo HTTP client to transfer a file to the local file system, modify your getFile class in order to be

This server fully supports HTTP ranges with no limitation related to the size of the requested object.

Java Tips

FileOutputStream fos = new FileOutputStream(f,true); // append mode

Methods seek(), length() and skipBytes() can be used to read or write slices of a file. Examples: